

Identifying Economic Regimes: *Reducing Downside Risks for University Endowments and Foundations*

JOHN M. MULVEY AND HAN LIU

JOHN M. MULVEY is a professor in the Operations Research and Financial Engineering Department and a founding member of the Bendheim Center for Finance at Princeton University in Princeton, NJ. mulvey@princeton.edu

HAN LIU is an assistant professor in the Operations Research and Financial Engineering Department and a member of the Statistics and Machine Learning Center at Princeton University in Princeton, NJ. hanliu@princeton.edu

The dynamic nature of financial markets presents a major challenge for long-term investors. One of the most durable, stylized facts involves contagion during severe economic conditions. Not only does the volatility of most securities increase, but also the correlation of returns can approach unity, causing securities to move together downward in a tight pack. The contagion environment has several causes—including plunging liquidity, a spike in systemic risk premium due to bank interlinkages and related issues, and panicking investors who often unload liquid securities to protect themselves against further losses.

Kyle and Xiong [2001] and Goldstein and Razin [2015] discuss alternative theories of financial crises. The empirical evidence during crash periods such as 2008 supports these suppositions. Importantly, for portfolio construction, contagion disrupts the anticipated diversification built into traditional models, and accordingly, the investor's portfolio performs much worse than would be expected. In turn, drawdown can be severe, leading to a major setback for long-term investors who depend upon capital to achieve their goals and pay future liabilities.

We build on the hypothesis that asset performance over long time periods can be separated into distinctive classes, called regimes, which display common characteristics. This supposition is not new; there have been

numerous studies and methods for addressing regime identification. For examples, see Hamilton and Susmel [1994], Timmermann [2000], Bilgili [2009], and Bae et al. [2013].

A multiregime analysis applies readily to nonprofit institutions, such as universities and foundations, for which the endowment supports a substantial portion of the organization's operating budget. For viability, it is critical that these institutions protect their endowment capital from severe erosion over long time horizons. Accordingly, there must be special attention to the downside risks. Institutions have several levers to improve their long-term viability. They can implement dynamic asset allocation strategies. They can adjust spending rules. And they might be able to seek additional funds from outside sources such as alumni and other benefactors.

There has been vast research on adjusting asset allocation based on changing economic conditions and regimes, such as van Norden and Schaller [1997], Ang and Bekaert [2002], Graflund and Nilsson [2003], Bauer et al. [2004], Guidolin and Timmermann [2007, 2008], Tu [2010], and Nystrup et al. [2015]. In contrast, the latter two levers are taken up herein.

The next section describes an algorithm from the area of machine learning to identify economic regimes. The method, called trend filtering, is one of the nonparametric approaches that have become common in the

machine-learning domain (Kim et al. [2009]). The trend-filtering algorithm categorizes regimes and in turn helps in the modeling of downside risk. The concepts are especially relevant for long-term investors who depend upon their endowment capital to provide a steady source of income.

To improve the worst-case outcomes, we show the advantages of adjustable spending rules during draw-down periods. As we will see, the two-regime approach assists in evaluating strategies to protect the long-term viability of an endowment for a nonprofit institution. The downside risks are closer to historical patterns than would be the case with a traditional static (single-regime) portfolio model.

APPLYING MACHINE LEARNING TO DETECT REGIMES

At its core, machine learning involves the classification of a “population” into classes possessing relatively homogeneous properties. Hastie et al. [2009] gave an excellent overview. In medicine, for example, there is intense interest in separating individuals into classes that have more or less likelihood of experiencing a disease such as coronary heart failure. To this point, apps and websites are available to help estimate the probability of a future heart attack (see the American Heart Association).¹ Herein, the questions depict “features” that support the underlying risk assessment. Given the goals of machine learning, it is natural that machine-learning methods may be applicable to detecting economic regimes for use in financial planning models.

Identifying economic regimes can be carried out in a variety of ways. A fundamental approach is to search for changes in the relationship between factors and the associated returns of assets via structural breaks, for example, Timmermann [2001]. Dating turning points between business cycles is an interesting and promising topic in economic research. In the United States, the National Bureau of Economic Research (NBER) calculates these turning points from historical data, which are widely used in business cycle analysis.

For our purposes, we focus on transition dates at which the economy alternates between a “growth” period and a “contraction” period. Several methods are applicable to real-time monitoring of the transition. Bry and Boschan [1971] established a turning point program, while Hamilton [1989] proposed a Markov switching model. Layton [1996] compared the two models.


Harding and Pagan [2003] added a Kalman filter approximation. Chauvet and Piger [2008] evaluated a dynamic factor time series model with the Harding and Pagan approach. Alternatively, a hidden Markov model (HMM) can help distinguish growth periods and crash periods in a probabilistic framework—for example, Cai [1994], Liew and Siu [2010], and Song [2014].

In our study, we aim to identify regimes of the financial markets in a data-driven and model-free context. The algorithm starts with the S&P 500 Index historical total return data. The basic idea is as follows. We treat the S&P 500 wealth path as a nonparametric function of time. If we naively plot this function, it will be a rough curve with a lot of noise. Therefore, we apply smoothing to the nonparametric function, which will “filter out” the noise of the historical data. The resulting function has a clear trend over a persistent period of time, which can help pinpoint regimes of the financial market. Compared to other approaches, this method is purely data-driven and does not rely on the assumption of a certain underlying probability model of the financial returns. Tibshirani [2014] and Liu et al. [2016] supply details about the approach and methods for addressing robustness from the machine-learning literature.

The trend-filtering algorithm identifies regimes by smoothing a piecewise linear function, labeling growth periods with 1 and contraction periods with -1 , with a regularization (penalty) term to dampen the number of changes in regime across a historical period—1990–2015 in the experiments. The algorithm discovers a reasonable compromise between stability in the curve and errors in the fit, as is common in the machine-learning domain (Tibshirani [1996, 2014] and Hastie et al. [2009]). The trend-filtering method deploys an optimization algorithm to directly determine the “status” of each historical time period. Exhibit 1 shows one of the main results.

EXHIBIT 1

Two Regimes for the S&P 500 Index, 1990–2015



It displays a two-regime plot of the S&P 500 Index over the 1990–2015 period, where light periods indicate growth and dark periods indicate contraction/crash.

Given the classification of the S&P 500 time series, we estimate the probabilities of movement from the state (growth or contraction) in period t to the state (growth or contraction) in the next period, $t + 1$. Exhibit 2 displays the equilibrium transition matrix across time periods (quarters). Thus, if the current quarter is in growth, there is a 90% probability that the next quarter will be in growth, and a 10% chance that the next quarter will be in contraction. Conversely, if the current quarter is in contraction, there is 60% probability of contraction in the following quarter. The transition probabilities form the basis of the switching procedure for the Monte Carlo simulations ahead. The equilibrium values are most

appropriate for long-term planning purposes, rather than short-term forecasting.

The next step in developing a long-term planning model is to select the asset categories. Exhibit 3 provides eight common assets and a typical allocation for large U. S. university endowments and foundations. Many of these institutions have followed the shift to alternative categories, as advocated by David Swensen [2009]. Liebowitz et al. [2010] discussed the endowment model. Mulvey and Holen [2016] surveyed current asset categories for the largest university endowments in the United States; the categories and associated allocation in Exhibit 3, Panel A, depict a representative portfolio of the larger endowments.

To build a forward-looking planning model, we must estimate future performance over the long planning horizon. To this end, there are a host of methods for estimating capital market assumptions and calibrating stochastic forecasting models. As illustrations, see Fabozzi et al. [2010], Ilmanen [2011], and Bogle and Nolan [2015]. For our case example, we surveyed multiple institutional sources and derived consensus projections in Exhibit 3, Panel B (adjusted for inflation). These values are relatively conservative compared with historical performance because of the current ultralow interest rates and earnings yield, and the modest growth of cyclically adjusted earnings (Pedersen [2015]).

EXHIBIT 2

Equilibrium Transition Matrix (probability, period t to $t + 1$)

	Growth Regime at Time $t + 1$	Contraction Regime at Time $t + 1$
Growth Regime at Time t	0.9	0.1
Contraction Regime at Time t	0.4	0.6

EXHIBIT 3

Summary Statistics for Multiperiod Simulation Study

	Private Equity	Real Estate	Hedge Fund	Real Assets	U.S. Equities	International Equity—Developed	International Equity—Emerging	U.S. Government Bond
Panel A: Asset Allocation								
Allocation	20.0%	12.0%	18.0%	7.0%	15.0%	12.0%	8.0%	8.0%
Panel B: Projected Returns for Asset Categories—Single-Regime (inflation adjusted)								
Annualized Rate								
Single-Regime Returns	6.500%	5.500%	5.000%	4.000%	4.500%	4.200%	4.800%	1.000%
Annual Volatility	14.212%	9.827%	7.681%	6.045%	17.988%	19.503%	26.344%	5.546%
Panel C: Historical Returns for Assets under Growth Regime (inflation adjusted)								
Annualized Rate								
Return under Growth	15.000%	9.000%	8.500%	4.500%	17.000%	16.500%	18.000%	0.300%
Annual Volatility	12.089%	7.881%	5.887%	6.157%	13.556%	15.731%	23.762%	4.236%
Panel D: Historical Returns for Assets under Contraction Regime (inflation adjusted)								
Annualized Rate								
Return under Crash	−21.665%	−7.412%	−7.907%	2.024%	−33.497%	−33.314%	−34.795%	3.849%
Annual Volatility	9.455%	12.345%	8.678%	5.561%	12.898%	13.654%	18.716%	7.601%

For simplicity, the volatility values reference historical performance.

When estimating the expected values under the two regimes, the relative performance of the historical values over the two regimes is preserved for the forward regimes. The goal is to generate similar “average” compound returns between the single-regime and the multiregime scenarios for the mean of all of the generated scenarios, while maintaining the relative performance across the regimes. The expected values in Exhibit 3, Panels C and D, are calculated so that this goal is achieved (shown in the simulation results).

Several items are worth mentioning. First, although the trend-filtering algorithm employs only the U.S. equity market in its regime identification, the returns of the other asset categories under the two regimes are consistent with a priori beliefs. For example, the international equity markets give performance similar to that of the U.S. markets. Likewise, U.S. government bond returns perform much better during contraction than in growth because of the flight-to-quality pattern during crash periods. The other assets display anticipated relationships as well. Because of the fundamental causes of contagion during contraction periods, we might expect similar patterns to occur in the future; we build on this assumption in our long-term planning model.

Applying Multiregime Simulations to Reduce Downside Risks

In this section, we implement Monte Carlo simulations for long-term investors. The case involves an institution that supports a substantial portion of its operating budget from endowment capital.² The endowment’s policy statement lists two main financial objectives: (1) provide steady income to maintain and improve the organization’s activities and (2) protect the endowment capital over long time periods. At times, the goals will conflict with each other; the trustees in conjunction with the endowment staff must set procedures to balance these goals.

As a background, prominent U.S. universities have established spending rules for reconciling these two goals. A typical rule might be to spend 3%–4% of the endowment capital each year for the operating budget and other needs, with the capital amount averaged over the previous several years. The 3%–4% range is deemed sufficiently conservative to achieve the second

goal—protecting capital over long time periods—whereas the averaging process reduces the volatility of spending so that large swings are avoided.

In this section, we evaluate spending rules in conjunction with a typical asset allocation by means of multiyear Monte Carlo simulation models. The experiments will compare a traditional single-regime model with a multiregime version. In the latter case, we will employ two regimes to illustrate the methodology. It is a simple matter to extend the approach to a larger number of regimes (Liu et al. [2016]).

The initial baseline simulations assume eight asset categories, an asset mix that is similar to many leading universities, and a single-regime structure (Exhibit 3, Panels A and B). In addition, the asset mix is set to the target allocation each period (quarter). Two extensions to the baseline are proposed. First, we implement the two-regime switching model with a focus on the downside risks (goal #2). Next, the analysis considers conditional variations on the spending rule, modifying spending when shocks occur. Herein, spending will decrease for a fixed period during and after a protracted drawdown episode.

The baseline experiments employ a single set of portfolio parameters. In particular, a fixed covariance matrix Σ of asset returns is derived from historical performance from 1990 to 2015, as well as an expected return vector $= \mu$ (Exhibit 3, Panel B). The expected returns are forward looking, based on a consensus forecast from several sources. Next, the returns are sampled from a multinormal distribution $N(\mu, \Sigma)$. Given these assumptions, we can readily generate random future paths for the endowment capital, along with associated spending values over the modeling horizon. Each path is designated as a scenario. The Monte Carlo simulation will evaluate a relatively large number of scenarios. For each random scenario s , the generation formula is

$$r_s = \mu + V * \varepsilon_s \quad (1)$$

where μ is the expected return vector, V is a Cholesky factor of Σ (i.e., $V V^t = \Sigma$), and ε_s is a random normal vector with zero mean and unit volatility $= N(0, I)$.

For simplicity, the return generating process assumes temporal independence. Thus, at each quarter t and under each scenario $s \in \mathcal{S}$, we generate a random vector of returns for the eight assets, based only on Equation 1. The simulation takes place over a long time

horizon, equal to 50 years in our case. The experiments employ 10,000 scenarios; each is a consistent path of random returns over the 2016 to 2066 time period for the eight assets. This type of forward simulation is relatively common for large institutional investors such as pension plans.

The baseline spending rule goes as follows. The target is to spend 4% of endowment capital each year: 0.5% from annual giving and other sources, and 3.5% from the endowment capital, with the capital averaged over the previous four years. For the first four years, the target is implemented. After year four, the target is compared with a specified range—3.5% to 4.75% of current capital. This range is enforced whenever the target falls outside, by setting spending at the top or bottom of the range.

The results of the baseline simulation appear in Exhibit 4. Exhibit 5 displays wealth paths for several quantiles across the 10,000 scenarios. The majority of scenarios indicate growth of the endowment capital. To analyze the worst cases, we fix three risk measures: (1) the probability that drawdown, in inflation-adjusted terms, will exceed 25% over the next 5 years, (2) the probability of 25% or worse drawdown over 10 years, and (3) the probability of a 50% or worse drawdown over 50 years.

The baseline risk numbers (Exhibit 4) appear reasonable, given that the institution generates positive cash flows at the 4% target to fund its operations—its primary goal. The 25% drawdown event probabilities

over 5- and 10-year horizons are approximately 10% and 20%, respectively; although significant, they are to be expected in an uncertain world fraught with economic and political risks. And the 4.9% probability of a 50% or worse drawdown event over the next 50 years is again within a reasonable range. Similarly, the mean values of the endowment continue to increase (starting at 1): 1.07 at 5 years, 1.15 after 10 years, and 1.58 after 50 years. The baseline spending rule and the asset allocation are relatively consistent with each other, and the trustees' responsibilities seem to fall within a prudent person criterion—at least to the authors' judgment. However, the simulation results depend upon the single-regime framework.

In this second set of experiments, we separate each quarter into one of two states: growth or contraction. The simulations begin at the first quarter of 2016 in the growth state. The state for the next period (quarter) will depend upon the equilibrium Markov transition matrix (Exhibit 2). There is a 90% chance that the next period will be growth; and if so, the random returns will be drawn from the multinormal distribution under growth. Alternatively, there is a 10% chance that the next quarter will be contraction; and in this case, the random returns for the eight assets will be drawn from the multinormal distribution under contraction. And so it goes.

A full path under the two-regime model will be a mixture driven by this transition matrix, with returns coming from the two multinormal distributions, sometimes growth and sometimes contraction. On average,

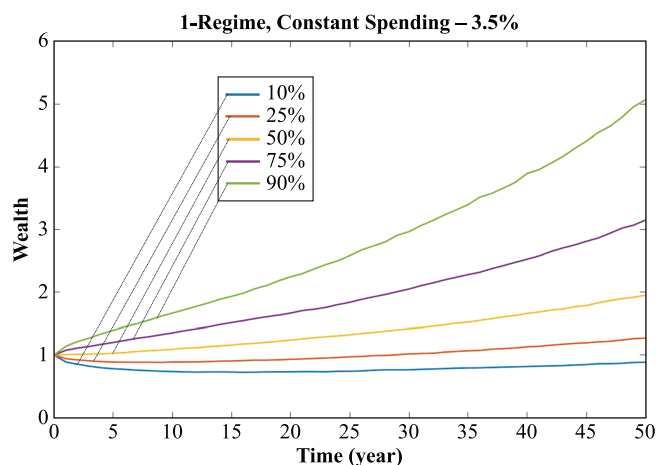
EXHIBIT 4

Summary Statistics for Baseline Monte Carlo Simulations (4% spending target)

Simulation Results	Without Spending-Cut Rule		Spending Cut by 20%	
	1-Regime	2-Regime	1-Regime	2-Regime
Crash Prob, 5 years	10.3%	18.4%	10.3%	18.4%
Crash Prob, 10 years	20.5%	31.8%	20.0%	31.5%
Crash Prob, 50 years	4.9%	19.9%	2.2%	13.1%
Mean-5 years	1.0644	1.0780	1.0644	1.0780
Mean-10 years	1.1635	1.1793	1.1692	1.1879
Mean-20 years	1.3998	1.4230	1.4173	1.4514
Mean-50 years	2.6630	2.6197	2.7152	2.7147
# of Simulations	10,000	10,000	10,000	10,000
Average % Time in "Adverse"	2.88%	5.60%	2.22%	4.60%

EXHIBIT 5

Wealth Paths across All Scenarios for Baseline



the number of quarters in growth will be about 80%, whereas the number of quarters in contraction will be about 20%, which is consistent with the historical record. The identical asset allocation (Exhibit 3, Panel A) is implemented over 10,000 scenarios, drawn with the same random number of seeds.

Exhibit 6 depicts an example of the future wealth path of the S&P 500 Index for a single scenario over the 50-year planning horizon. Growth and contraction periods are indicated. The simulation generates 10,000 such paths.

Turning to the initial two-regime results (Exhibit 4), we see that the size of the mean endowment over the 10,000 scenarios is comparable between the single and multiregime simulations at each future time juncture (5, 10, 20, and 50 years). Also, on average, the endowment continues to grow under two regimes. However, the probabilities of worst-case outcomes are much higher; for example, there is a 19.9% probability of 50% drawdown over 50 years under two regimes, compared with 4.9% under the single regime. A similar relationship occurs for the other two risk measures. Clearly, modeling contagion in the simulation causes the worst-case outcomes to be higher.

Exhibit 7 provides another view of the left tails by plotting the respective cumulative distributions; the downside risks are much lower for the single-regime simulation compared with the two-regime model. Exhibit 8 displays the ranges across scenarios for the two-regime simulation, again showing the wider ranges at the lower extremes for the two-regime model (compare Exhibits 5 and 8).

Do the new higher risks exceed the reasonableness threshold? The guardians of the institution, typically the trustees, have the primary responsibility of addressing this question. A more appropriate question

for the reader is, what can be done to reduce the downside risks? Although there are a number of options, an obvious area to explore involves institutional spending. One approach is to manage the drawdown periods by adjustments to spending under adverse circumstances. And the smoothing of the endowment capital over a multiyear back period is a small step in this direction. Another approach is to reduce spending in a dynamic fashion when large drawdown thresholds are breached. This dynamic spending rule is analogous to peak pricing strategies in energy applications.

To illustrate the impact of conditional spending on the worst-case outcomes, the next Monte Carlo

EXHIBIT 7

Cumulative Distribution of Wealth at 50-Year Horizon (inflation adjusted)




EXHIBIT 8

Wealth Paths Across All Scenarios—Two-Regime Model

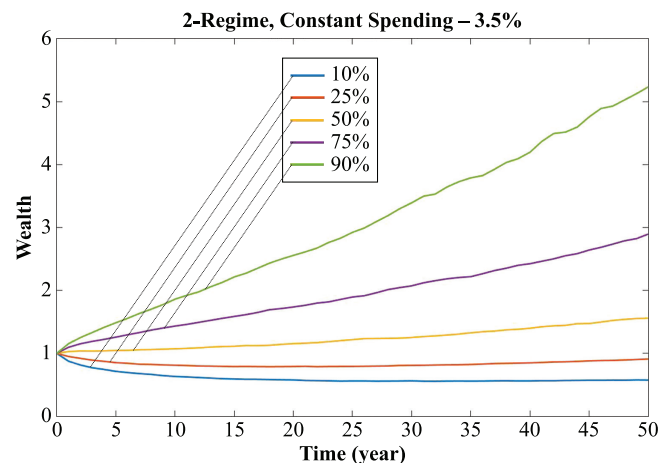



EXHIBIT 6

A Representative Scenario Path over the 50-Year Planning Horizon (S&P 500 Index)



simulations lower spending during and after a significant drawdown. Spending drops by 20% for three years when drawdown exceeds 20%, as measured by the initial capital level. Exhibit 4 shows the impact of this dynamic spending rule. As expected, the reduction has little impact over the 5- and 10-year horizons, because of the rarity of entering the so-called adverse periods. On the other hand, dynamic spending has a noticeable improvement in the worst-case risks over the 50-year period, for instance, dropping from 19.9% to 13.1% under the two-regime case. Note, also, that the total time in adverse events drops as well. Reducing spending at one period helps protect against adverse events in later periods.

Importantly, changes to the base spending (4% inflation adjusted per year) can be part of a sensitivity exercise. To wit, Exhibit 9 depicts the impacts. Increasing the spending rate to a 4.5% target per year has a noticeable impact on the worst-case outcomes. Not surprisingly, decreasing the spending rate to 3.5% per year

reduces the downside risks. A conservative strategy is to spend at the 3.5% target (0.5% giving and 3% from endowment capital) per year, coupled with the dynamic spending rule—resulting in the lowest probability, 6.3%, of the 50% drawdown over the next 50 years. Again, the institution must select an acceptable trade-off among its opportunities, goals, and risks.

FUTURE DIRECTIONS

The multiregime framework gives a more realistic depiction of the left-tail events than the traditional static (single-regime) approach and thus provides better information to “simulate” the institution’s cash needs and viability in the future. In this way, the regime model helps the institution achieve its dual long-term goals of providing a steady source of income and protecting the endowment capital. To this end, the advantages of dynamic spending rules are quantified.

A primary purpose of this article has been to illustrate the importance of modeling contagion on the worst-case outcomes for long-term investors. Accordingly, we kept the modeling details sparse. Several extensions of the multiregime approach stand out.³ The two-regime approach can be extended to three regimes (growth, transition, and contraction), for example, as shown in Exhibit 10. Also, the Monte Carlo simulations can include significant considerations such as limitations on illiquid asset transactions, market impact costs, and other realistic constraints. A dynamic asset allocation strategy might be applied in a modest way, given the difficulty of moving across the illiquid alternative asset categories.

Likewise, the asset allocation decisions and the spending rules can be linked in an integrated fashion via


EXHIBIT 9

Summary Statistics for Multiperiod Simulation Study

Simulation Results	Without Spending Cut Rule		Spending Cut by 20%	
	1-Regime	2-Regime	1-Regime	2-Regime
Panel A: Performance Statistics for 4.5% Spending Target				
Crash Prob, 5 years	12.7%	20.8%	12.7%	20.8%
Crash Prob, 10 years	28.6%	38.2%	27.7%	37.7%
Crash Prob, 50 years	27.9%	43.3%	9.1%	27.3%
Mean-5 years	1.0324	1.0462	1.0324	1.0462
Mean-10 years	1.0702	1.0913	1.0806	1.1037
Mean-20 years	1.1510	1.1897	1.2021	1.2459
Mean-50 years	1.5819	1.6233	1.7786	1.8345
# of Simulations	10,000	10,000	10,000	10,000
Average % Time in “Adverse”	7.56%	9.51%	5.01%	7.32%
Panel B: Performance Statistics for 3.5% Spending Target				
Crash Prob, 5 years	8.3%	16.3%	8.3%	16.3%
Crash Prob, 10 years	14.6%	26.6%	14.4%	26.4%
Crash Prob, 50 years	1.2%	8.5%	0.4%	5.6%
Mean-5 years	1.0977	1.1112	1.0977	1.1112
Mean-10 years	1.2421	1.2609	1.2462	1.2680
Mean-20 years	1.5883	1.6267	1.6018	1.6499
Mean-50 years	3.6169	3.6499	3.6585	3.7310
# of Simulations	10,000	10,000	10,000	10,000
Average % Time in “Adverse”	1.39%	3.41%	1.06%	2.68%

EXHIBIT 10

Applying Trend Filtering with Three Regimes (growth, transition, contraction)



an “enterprise risk management” system for a nonprofit institution. Herein, the portfolio is expanded to include the endowment and the sponsoring organization. The reduced spending rule discussed in this report is only one version among others. Undoubtedly, there are better procedures for balancing the institution’s needs to deploy its endowment resources to achieve current and intermediate goals, while at the same time, protecting its capital over long time horizons.

As an example, in conjunction with the reduced spending, a nonprofit institution should look to implement operational strategies that can be reversed relatively easily during severe drawdown periods. For example, a recently proposed increase in sabbatical leaves for faculty at Princeton University, costing about \$10 to \$15 million per year (rough estimate by Mulvey), might be suspended during a crash. Other reductions might be possible: temporarily freezing wages, postponing the replacement of nonessential employees who have retired or resigned, and lagging noncritical capital improvements. Similarly, the institution may have some control over donations and contribution campaigns. As we see from the two-regime simulation, the long-term benefits of return-contingent spending rates illustrate how important it is for institutions to plan for such contingencies, despite how rarely they occur.⁴

ENDNOTES

We greatly appreciate the assistance of our doctoral students at Princeton University: Han Hao, Jing Ye, and Tianqi Zhao.

¹American Heart Association, “Heart Attack Risk Calculator,” www.heart.org, 2016.

²The case study involves a major university endowment, with inflation measured by the Higher Education Price Index (HEPI), which is somewhat higher than the CPI.

³Other tail-risk measures such as conditional value at risk can be readily evaluated within the multiregime framework.

⁴Reduced spending rules are pertinent for individuals during their retirement years. There are similar advantages for applying contingent spending rules to protect the retirement capital during severe drawdown periods.

REFERENCES

- Ang, A., and G. Bekaert. “International Asset Allocation with Regime Shifts.” *Review of Financial Studies*, Vol. 15, No. 4 (2002), pp. 1137–1187.
- Bae, G., W.C. Kim, and J. Mulvey. “Dynamic Asset Allocation for Varied Financial Markets under Regime Switching Framework.” *European Journal of Operational Research*, Vol. 234, No. 2 (2013), pp. 450–458.
- Bauer, R., R. Haerden, and R. Molenaar. “Asset Allocation in Stable and Unstable Times.” *The Journal of Investing*, Vol. 13, No. 3 (2004), pp. 72–80.
- Bilgili, M. “Clustering Techniques and Multi Regime Stochastic Optimization with Applications in Finance.” PhD dissertation, Princeton University, 2009.
- Bogle, J., and M. Nolan Jr. “Occam’s Razor Redux: Establishing Reasonable Expectations for Financial Market Returns.” *The Journal of Portfolio Management*, Vol. 42, No. 1 (2015), pp. 119–134.
- Bry, G., and C. Boschan. *Cyclical Analysis of Time Series: Selected Procedures and Computer Programs*. New York, NY: Columbia University Press, 1971.
- Cai, J. “A Markov Model of Unconditional Variance in ARCH.” *Journal of Business and Economic Statistics*, Vol. 12, No. 1 (1994), pp. 309–316.
- Chauvet, M., and J. Piger. “A Comparison of the Real-Time Performance of Business Cycle Dating Methods.” *Journal of Business and Economic Statistics*, Vol. 26, No. 1 (2008), pp. 42–49.
- Fabozzi, F., S. Focardi, and P. Kolm. *Quantitative Equity Investing: Techniques and Strategies* (Frank J. Fabozzi Series). Hoboken, NJ: John Wiley & Sons, 2010.
- Goldstein, I., and A. Razin. “Three Branches of Theories of Financial Crises.” *Foundations and Trends in Finance*, Vol. 10, No. 2 (2015), pp. 113–180.
- Graflund, A., and B. Nilsson. “Dynamic Portfolio Selection: the Relevance of Switching Regimes and Investment Horizon.” *European Financial Management*, Vol. 9, No. 2 (2003), pp. 179–200.
- Guidolin, M., and A. Timmermann. “Asset Allocation Under Multivariate Regime Switching.” *Journal of Economic Dynamics and Control*, Vol. 31, No. 11 (2007), pp. 3503–3544.
- . “International Asset Allocation under Regime Switching, Skew, and Kurtosis Preferences.” *Review of Financial Studies*, Vol. 21, No. 2 (2008), pp. 889–935.

- Hamilton, J.D. "A New Approach to the Economic Analysis of Nonstationary Time Series and the Business Cycle." *Econometrica*, Vol. 57, No. 2 (1989), pp. 357-384.
- Hamilton, J.D., and R. Susmel. "Autoregressive Conditional Heteroskedasticity and Changes in Regime." *Journal of Econometrics*, Vol. 64, No. 1 (1994), pp. 307-333.
- Harding, D., and A. Pagan. "A Comparison of Two Business Cycle Dating Methods." *Journal of Economic Dynamics and Control*, Vol. 27, No. 9 (2003), pp. 1681-1690.
- Hastie, T., R. Tibshirani, and J. Friedman. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, 2nd ed. New York, NY: Springer, 2009.
- Ilmanen, A. *Expected Returns: An Investor's Guide to Harvesting Market Rewards*. Chichester, U.K.: John Wiley & Sons, 2011.
- Kim, S.-J., K. Koh, S. Boyd, and D. Gorinevsky. "L1 Trend Filtering." *SIAM Review*, Vol. 51, No. 2 (2009), pp. 339-360.
- Kyle, A., and W. Xiong. "Contagion as a Wealth Effect." *Journal of Finance*, 56 (2001), pp. 1401-1440.
- Layton, A.P. "Dating and Predicting Phase Changes in the U.S. Business Cycle." *International Journal of Forecasting*, Vol. 12, No. 3 (1996), pp. 417-428.
- Liebowitz, M., A. Bova, and P.B. Hammond. *The Endowment Model of Investing: Return, Risk, and Diversification*. Hoboken, NJ: John Wiley & Sons, 2010.
- Liew, C.C., and T.K. Siu. "A Hidden Markov Regime-Switching Model for Option Valuation." *Insurance Mathematics and Economics*, Vol. 47, No. 3 (2010), pp. 374-384.
- Liu, H., J. Mulvey, and T. Zhao. "A Nonparametric Smoothing Approach to Financial Market Regime Identification." Princeton University Report, February 2016.
- Mulvey, J., and M. Holen. "What Can We Learn from University Endowments: Insights from Asset Categories." *Journal of Investment Consulting*, forthcoming 2016.
- Nystrup, P., B.W. Hansen, H. Madsen, and E. Lindstrom. "Regime-Based Versus Static Asset Allocation: Letting the Data Speak." *The Journal of Portfolio Management*, Vol. 42, No. 1 (2015), pp. 103-109.
- Pedersen, L.H. *Efficiently Inefficient: How Smart Money Invests and Market Prices Are Determined*. Princeton NJ: Princeton University Press, 2015.
- Song, Y. "Modelling Regime Switching and Structural Breaks with an Infinite Hidden Markov Model." *Journal of Applied Econometrics*, Vol. 29, No. 5 (2014), pp. 825-842.
- Swensen, D. *Pioneering Portfolio Management: An Unconventional Approach to Institutional Investment*, 2nd ed. New York: Free Press, 2009.
- Tibshirani, R. "Regression Shrinkage and Selection via the Lasso." *Journal of the Royal Statistical Society. Series B (Methodological)*, Vol. 58, No. 1 (1996), pp. 267-288.
- . "Adaptive Piecewise Polynomial Estimation via Trend Filtering." *Annals of Statistics*, Vol. 42, No. 1 (2014), pp. 285-323.
- Timmermann, A. "Moments of Markov Switching Models." *Journal of Econometrics*, Vol. 96, No. 1 (2000), pp. 75-111.
- . "Structural Breaks, Incomplete Information, and Stock Prices." *Journal of Business & Economic Statistics*, Vol. 19, No. 3 (2001), pp. 299-314.
- Tu, J. "Is Regime Switching in Stock Returns Important in Portfolio Decisions?" *Management Science*, Vol. 56, No. 7 (2010), pp. 1198-1215.
- van Norden, S., and H. Schaller. "Regime Switching in Stock Market Returns." *Applied Financial Economics*, Vol. 7, No. 2 (1997), pp. 177-191.
- To order reprints of this article, please contact Dewey Palmieri at dpalmieri@ijournals.com or 212-224-3675.*